

---

# Hive

---

OutLine

Hive基础

【实践】Hive搭建

【实践】Hive练习

## 背景

- 引入原因：
  - 对存在HDFS上的文件或HBase中的表进行查询时，是要手工写一堆MapReduce代码
  - 对于统计任务，只能由动MapReduce的程序员才能搞定
  - 耗时耗力，更多精力没有有效的释放出来
- **Hive基于一个统一的查询分析层，通过SQL语句的方式对HDFS上的数据进行查询、统计和分析**

## Hive 是什么

- Hive是一个SQL解析引擎，将SQL语句转译成MR Job，然后再Hadoop平台上运行，达到快速开发的目的。
- Hive中的表是纯逻辑表，就只是表的定义等，即表的元数据。本质就是Hadoop的目录/文件，达到了元数据与数据存储分离的目的
- Hive本身不存储数据，它完全依赖HDFS和MapReduce。
- Hive的内容是读多写少，不支持对数据的改写和删除
- Hive中没有定义专门的数据格式，由用户指定，需要指定三个属性：
  - 列分隔符
  - 行分隔符
  - 读取文件数据的方法

```

* Licensed to the Apache Software Foundation (ASF) under one
* or more contributor license agreements. See the NOTICE file
* distributed with this work for additional information
* regarding copyright ownership. The ASF licenses this file
* to you under the Apache License, Version 2.0 (the
* "License"); you may not use this file except in compliance
* with the License. You may obtain a copy of the License at
* http://www.apache.org/licenses/LICENSE-2.0
* Unless required by applicable law or agreed to in writing,
* software distributed under the License is distributed on an
* "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY
* KIND, either express or implied. See the License for the
* specific language governing permissions and limitations under
* the License.
*/
/**
 * This is an example Hadoop Map/Reduce application.
 * It reads the input, counts the number of occurrences of each word,
 * and counts the total number of words.
 * A reducer class that just emits the sum of the input values.
 */
public static class WordCount implements Reducer<Text, IntWritable> {
    /**
     * To run: bin/hadoop jar wordcount.jar WordCount
     */
    public void reduce(Text key, Iterable<IntWritable> values,
        OutputCollector<IntWritable> collector, Reporter reporter)
        throws IOException {
        int sum = 0;
        while (values.iterator().hasNext()) {
            sum += values.iterator().next().get();
        }
        collector.collect(new IntWritable(sum));
    }
}

public class WordCountDriver {
    /**
     * The main driver for word count map/reduce program.
     * Invoke this method to submit the map/reduce job.
     * @throws IOException When there is communication problems with the
     * job tracker.
     */
    public int run(String[] args) throws Exception {
        JobConf conf = new JobConf(getConf(), WordCount.class);
        conf.setJobName("wordcount");

        // the keys are words (strings)
        conf.setOutputKeyClass(Text.class);
        // the values are counts (ints)
        conf.setOutputValueClass(IntWritable.class);

        conf.setMapperClass(MapClass.class);
        conf.setCombinerClass(Reduce.class);
        conf.setReducerClass(Reduce.class);

        List<String> other_args = new ArrayList<String>();
        for(int i=0; i < args.length; ++i) {
            try {
                if ("m".equals(args[i])) {
                    conf.setNumMapTasks(Integer.parseInt(args[++i]));
                } else if ("r".equals(args[i])) {
                    conf.setNumReduceTasks(Integer.parseInt(args[++i]));
                } else {
                    other_args.add(args[i]);
                }
            } catch (NumberFormatException except) {
                System.out.println("ERROR: Integer expected instead of " + args[i]);
                return printUsage();
            } catch (ArrayIndexOutOfBoundsException except) {
                System.out.println("ERROR: Required parameter missing from " +
                    args[i-1]);
                return printUsage();
            }
        }
        // Make sure there are exactly 2 parameters left.
        if (other_args.size() != 2) {
            System.out.println("ERROR: Wrong number of parameters: " +
                other_args.size() + " instead of 2.");
            return printUsage();
        }
        FileInputFormat.setInputPaths(conf, other_args.get(0));
        FileOutputFormat.setOutputPath(conf, new Path(other_args.get(1)));

        JobClient.runJob(conf);
        return 0;
    }
}

```

```
select word, count(*)
from (
select
explode(split(sentence, ' ')) as word
from article
) t
group by word
```

# 简单!!!

部资料，盗版必究——

# Hive 中的 SQL 与传统 SQL 区别

	HQL	SQL
数据存储	HDFS、Hbase	Local FS
数据格式	用户自定义	系统决定
数据更新	不支持（把之前的数据覆盖）	支持
索引	有（0.8版之后增加）	有
执行	MapReduce	Executor
执行延迟	高	低
可扩展性	高（UDF、UDAF, UDTF）	低
数据规模	大（数据大于TB）	小
数据检查	读时模式	写时模式

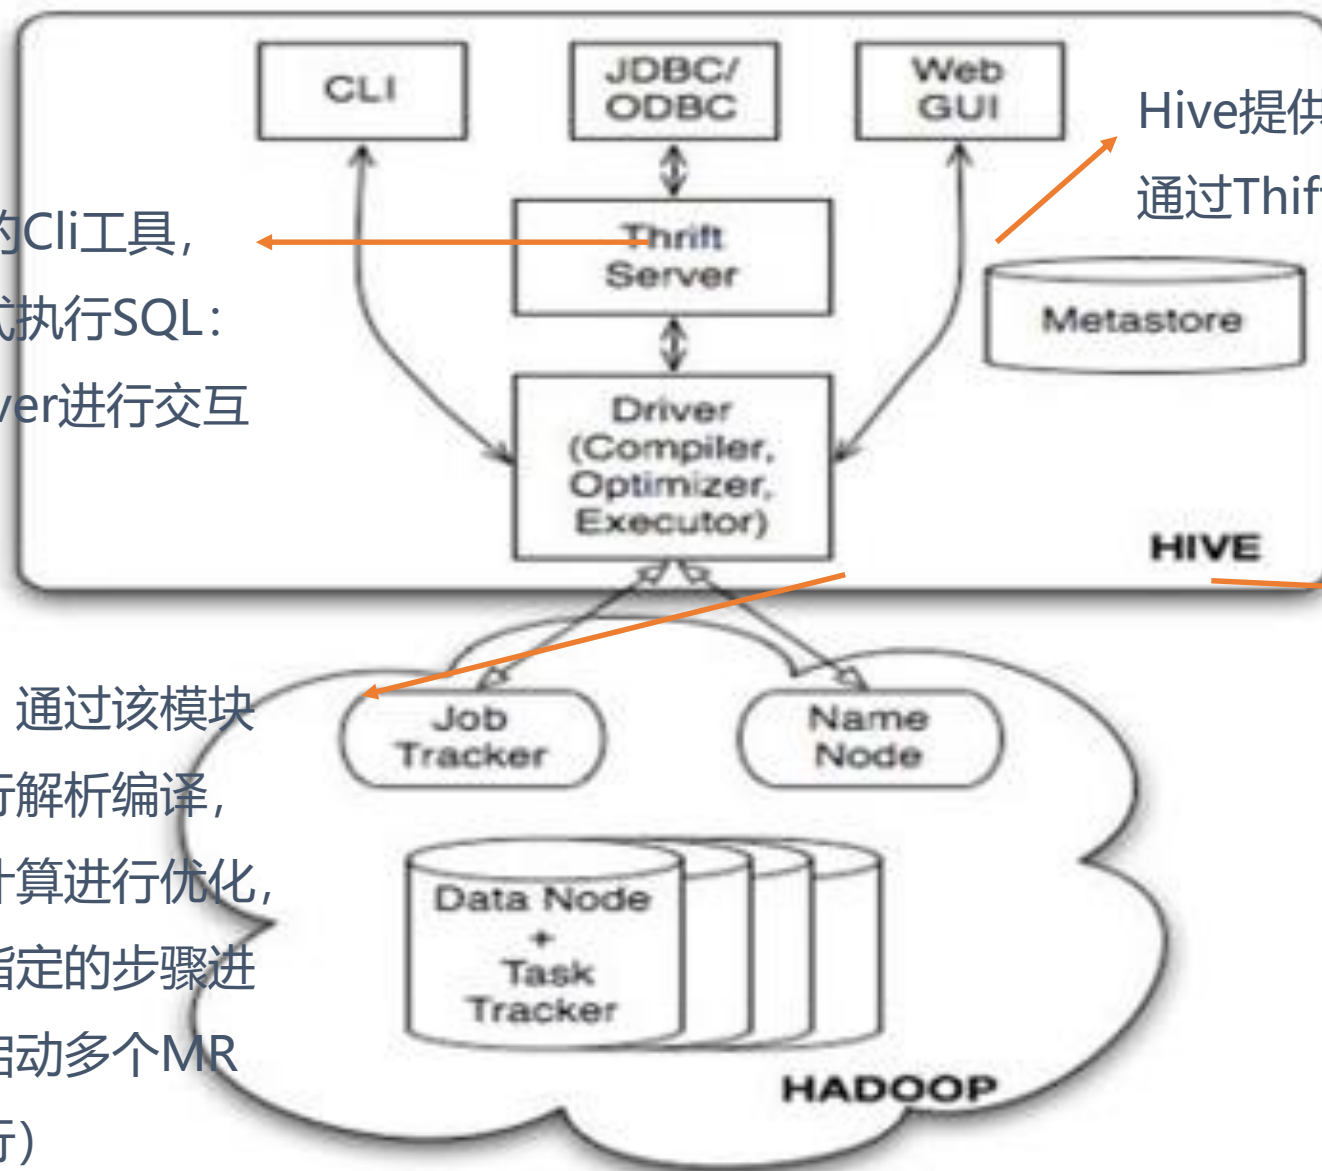
## 与传统关系数据特点比较

- hive和关系数据库存储文件的系统不同，hive使用的是hadoop的HDFS（hadoop的分布式文件系统），关系数据库则是服务器本地的文件系统；
- hive使用的计算模型是mapreduce，而关系数据库则是自己设计的计算模型；
- 关系数据库都是为实时查询的业务进行设计的，而hive则是为海量数据做数据挖掘设计的，实时性很差
- Hive很容易扩展自己的存储能力和计算能力，这个是继承hadoop的，而关系数据库在这个方面要比数据库差很多。

## Hive 体系架构

Hive提供的Cli工具，  
进行交互式执行SQL：  
直接与Driver进行交互

驱动模块：通过该模块  
对输入进行解析编译，  
对需求的计算进行优化，  
然后按照指定的步骤进  
行（通常启动多个MR  
任务来执行）




Hive提供JDBC驱动，作为JAVA的API：JDBC是  
通过Thrift Server来接入，然后发送给Driver

元数据：是一个独立的关系型数  
据库，通常Mysql，Hive会在其  
中保存表模式和其他系统元数据



## Hive 体系架构



## 用户接口

- CLI: 启动的时候, 会同时启动一个 Hive 副本
- Client: Hive 的客户端, 用户连接至 Hive Server
- GUI: 通过浏览器访问 Hive

## 语句转换

- 解析器: 生成抽象语法树
- 语法分析器: 验证查询语句
- 逻辑计划生成器 (包括优化器): 生成操作符树
- 查询计划生成器: 转换为map-reduce任务

## 数据存储

- Hive数据以文件形式存储在HDFS的指定目录下
- Hive语句生成查询计划, 由MapReduce调用执行

## Hive 数据管理

- hive的表本质就是Hadoop的目录/文件
  - hive默认表存放路径一般都是在你的工作目录的hive目录里面，按表名做文件夹分开，如果你有分区表的话，分区值是子文件夹，可以直接在其它的M/R job里直接应用这部分数据

	Name	HDFS Directory
Table	mobile_user	/lbs/mobile_user
Partition	action = insight, day= 20131020	/lbs/mobile_user/action=insight/day=20131020
Bucket	clusted by user into 32 buckets	/lbs/mobile_user/action=insight/day=20131020/part-00000

## Hive 内部表和外部表

- Hive的create创建表的时候，选择的创建方式：
  - create table
  - create external table
- 特点：
  - 在导入数据到外部表，数据并没有移动到自己的数据仓库目录下，也就是说外部表中的数据并不是由它自己来管理的！而表则不一样；
  - 在删除表的时候，Hive将会把属于表的元数据和数据全部删掉；而删除外部表的时候，Hive仅仅删除外部表的元数据，数据是不会删除的！

## Hive 中的 Partition

- 在 Hive 中，表中的一个 Partition 对应于表下的一个目录，所有的 Partition 的数据都存储在对应的目录中
  - 例如：pvs 表中包含 ds 和 city 两个 Partition，则
  - 对应于 ds = 20090801, ctry = US 的 HDFS 子目录为：/wh/pvs/ds=20090801/ctry=US;
  - 对应于 ds = 20090801, ctry = CA 的 HDFS 子目录为：/wh/pvs/ds=20090801/ctry=CA
- partition是辅助查询，缩小查询范围，加快数据的检索速度和对数据按照一定的规格和条件进行管理。

## Hive 中的 Bucket

- hive中table可以拆分成partition, table和partition可以通过 'CLUSTERED BY' 进一步分bucket, bucket中的数据可以通过 'SORT BY' 排序。
- create table bucket\_user (id int,name string)clustered by (id) into 4 buckets;
- 'set hive.enforce.bucketing = true' 可以自动控制上一轮reduce的数量从而适配bucket的个数, 当然, 用户也可以自主设置mapred.reduce.tasks去适配bucket个数
- Bucket主要作用:
  - 数据sampling

—— 八斗大数据内部资料, 盗版必究 ——

提升某些查询操作效率, 例如mapside join

## Hive 中的 Bucket

- 查看sampling数据：
  - hive> select \* from student tablesample(bucket 1 out of 2 on id);
  - tablesample是抽样语句，语法：TABLESAMPLE(BUCKET x OUT OF y)
  - y必须是table总bucket数的倍数或者因子。hive根据y的大小，决定抽样的比例。例如，table总共分了64份，当y=32时，抽取 $(64/32=)$ 2个bucket的数据，当y=128时，抽取 $(64/128=)$ 1/2个bucket的数据。x表示从哪个bucket开始抽取。例如，table总bucket数为32，tablesample(bucket 3 out of 16)，表示总共抽取  $(32/16=)$  2个bucket的数据，分别为第3个bucket和第  $(3+16=)$  19个bucket的数据。

## Hive 数据类型

- 数据类型

- 原生类型

- TINYINT
    - SMALLINT
    - INT
    - BIGINT
    - BOOLEAN
    - FLOAT
    - DOUBLE
    - STRING
    - BINARY (Hive 0.8.0以上才可用)
    - TIMESTAMP (Hive 0.8.0以上才可用)

## Hive 数据类型

- 数据类型

- 复合类型

- Arrays: ARRAY<data\_type>
    - Maps: MAP<primitive\_type, data\_type>
    - Structs: STRUCT<col\_name: data\_type[COMMENT col\_comment],.....>
    - Union: UNIONTYPE<data\_type, data\_type,.....>



## Hive SQL —— Join in MR

```
INSERT OVERWRITE TABLE pv_users  
SELECT pv.pageid, u.age  
FROM page_view pv  
JOIN user u  
ON (pv.userid = u.userid);
```

*page\_view*

pageid	userid	time
1	111	9:08:01
2	111	9:08:13
1	222	9:08:14

*user*

userid	age	gender
111	25	female
222	32	male

Map

key	value
111	<1,1>
111	<1,2>
222	<1,1>

Shuffle  
Sort

key	value
111	<2,25>
222	<2,32>

key	value
111	<1,1>
111	<1,2>
111	<2,25>

*pv\_users*

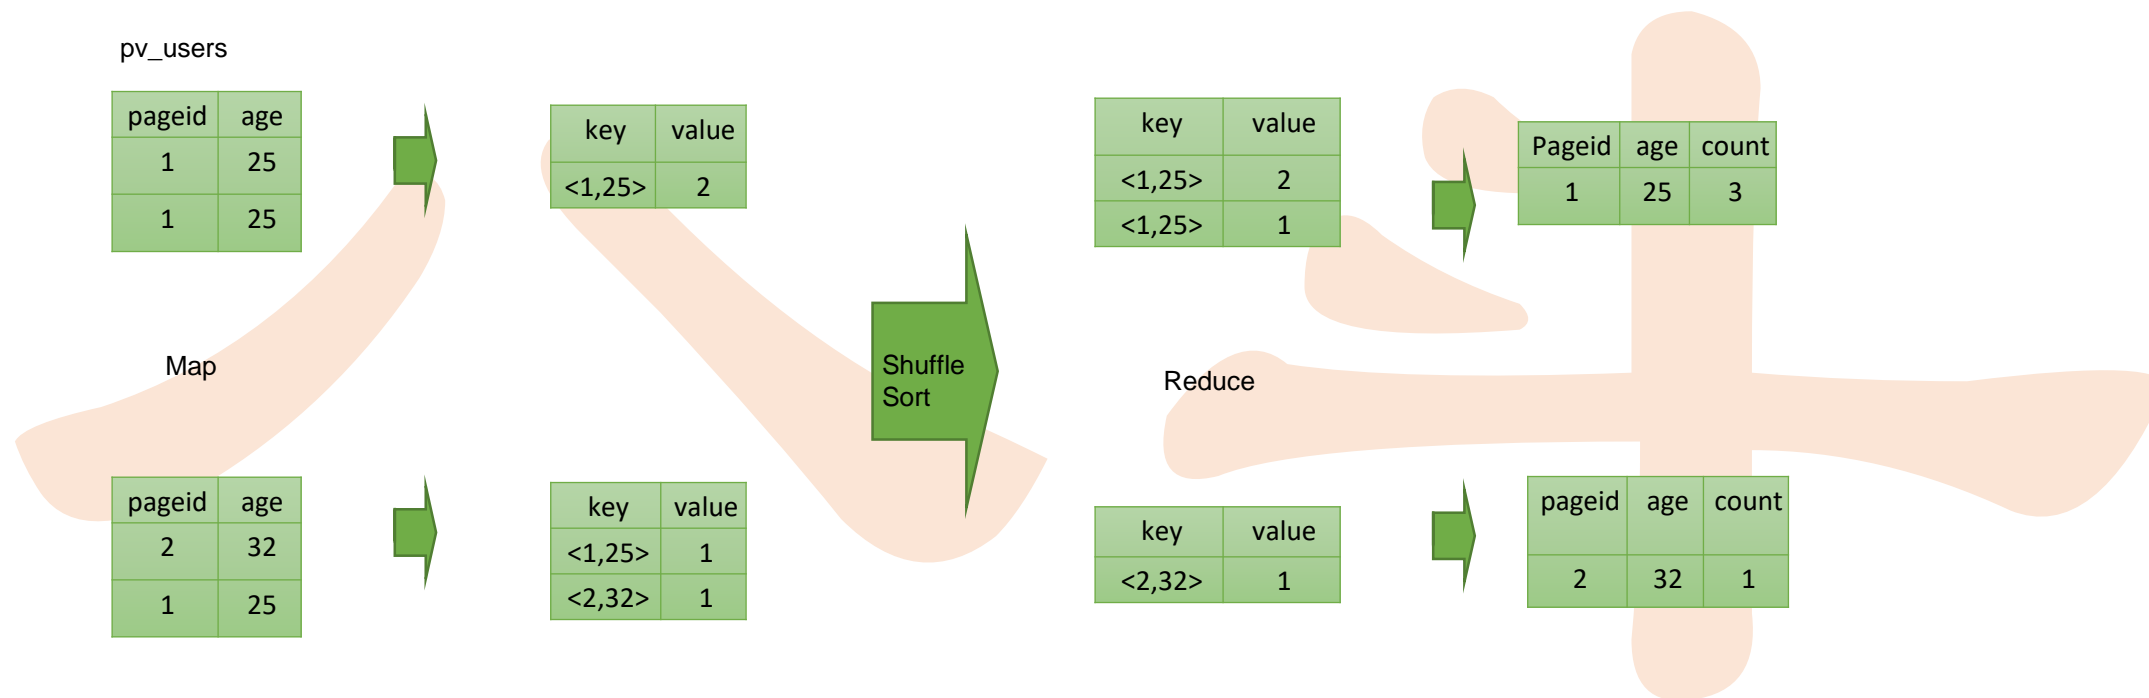
Pageid	age
1	25
2	25

Reduce

pageid	age
1	32

## Hive SQL —— Join in MR

***SELECT pageid, age, count(1)  
FROM pv\_users  
GROUP BY pageid, age;***



## Hive 的优化

- Map的优化:
  - 作业会通过input的目录产生一个或者多个map任务。set dfs.block.size(=128)
  - Map越多越好吗? 是不是保证每个map处理接近文件块的大小?
  - 如何合并小文件, 减少map数?

```
set mapred.max.split.size=100000000;  
set mapred.min.split.size.per.node=100000000;  
set mapred.min.split.size.per.rack=100000000;  
set hive.input.format=org.apache.hadoop.hive.ql.io.CombineHiveInputFormat;
```

- 如何适当的增加map数?

```
set mapred.map.tasks=10;
```

- Map端聚合 hive.map.aggr=true。Mr中的Combiners.

## Hive 的优化

- Reduce的优化:
  - `hive.exec.reducers.bytes.per.reducer`; reduce任务处理的数据量
  - 调整reduce的个数:
    - 设置reduce处理的数据量
    - `set mapred.reduce.tasks=10`
- 一个Reduce:
  - 没有group by
  - order by (可以使用distribute by和sort by)
  - 笛卡尔积

```
select pt,count(1)
from popt_tbaccountcopy_mes
where pt = '2012-07-04' group by pt;
写成
select count(1)
from popt_tbaccountcopy_mes
where pt = '2012-07-04';
```

## Hive 的优化

- 分区裁剪 (partition)
  - Where中的分区条件, 会提前生效, 不必特意做子查询, 直接Join和GroupBy
- 笛卡尔积
  - join的时候不加on条件或者无效的on条件, Hive只能使用1个reducer来完成笛卡尔积
- Map join
  - /\*+ MAPJOIN(tablelist) \*/ , 必须是小表, 不要超过1G, 或者50万条记录
- Union all
  - 先做union all再做join或group by等操作可以有效减少MR过程, 尽管是多个Select, 最终只有一个mr

## Hive 的优化

- Multi-insert & multi-group by
  - 从一份基础表中按照不同的维度，一次组合出不同的数据
  - FROM from\_statement
  - INSERT OVERWRITE TABLE tablename1 [PARTITION (partcol1=val1)] select\_statement1 group by key1
  - INSERT OVERWRITE TABLE tablename2 [PARTITION(partcol2=val2 )] select\_statement2 group by key2
- Automatic merge
  - 当文件大小比阈值小时，hive会启动一个mr进行合并
  - hive.merge.mapfiles = true 是否和并 Map 输出文件，默认为 True
  - hive.merge.mapredfiles = false 是否合并 Reduce 输出文件，默认为 False
  - hive.merge.size.per.task = 256\*1000\*1000 合并文件的大小
- Multi-Count Distinct
  - 必须设置参数：set hive.groupby.skewindata=true;
  - select dt, count(distinct uniq\_id), count(distinct ip)
  - from ods\_log where dt=20170301 group by dt

## Hive 的 Join 优化

- 一个MR job

```
SELECT a.val, b.val, c.val  
FROM a  
JOIN b ON (a.key = b.key1)  
JOIN c ON (a.key = c.key1)
```

- 生成多个MR job

```
SELECT a.val, b.val, c.val  
FROM a  
JOIN b ON (a.key = b.key1)  
JOIN c ON (c.key = b.key1)
```

## Hive 的 Join 优化 —— 表连接顺序

- 按照JOIN顺序中的最后一个表应该尽量是大表，因为JOIN前一阶段生成的数据会存在于Reducer的buffer中，通过stream最后面的表，直接从Reducer的buffer中读取已经缓冲的中间结果数据（这个中间结果数据可能是JOIN顺序中，前面表连接的结果的Key，数据量相对较小，内存开销就小），这样，与后面的大表进行连接时，只需要从buffer中读取缓存的Key，与大表中的指定Key进行连接，速度会更快，也可能避免内存缓冲区溢出。

```
SELECT /*+ STREAMTABLE(a) */ a.val, b.val, c.val  
FROM a  
JOIN b ON (a.key = b.key1)  
JOIN c ON (c.key = b.key1);  
a表被视为大表
```

```
SELECT /*+ MAPJOIN(b) */ a.key, a.value  
FROM a  
JOIN b ON a.key = b.key;
```

MAPJOIN会把小表全部读入内存中，在map阶段直接拿另外一个表的数据和内存中表数据做匹配，由于在map是进行了join操作，省去了reduce运行的效率也会高很多。



## Hive 的 Join 优化 —— 表连接顺序

- 左连接时，左表中出现的JOIN字段都保留，右表没有连接上的都为空。


```
SELECT a.val, b.val  
FROM a  
LEFT OUTER JOIN b ON (a.key=b.key)  
WHERE a.ds='2009-07-07' AND b.ds='2009-07-07'
```

```
ASELECT a.val, b.val  
FROM a  
LEFT OUTER JOIN b  
ON (a.key=b.key AND b.ds='2009-07-07' AND a.ds='2009-07-07')
```

- 执行顺序是，首先完成2表JOIN，然后再通过WHERE条件进行过滤，这样在JOIN过程中可能会输出大量结果，再对这些结果进行过滤，比较耗时。可以进行优化，将WHERE条件放在ON后，在JOIN的过程中，就对不满足条件的记录进行了预先过滤。

## Hive 的优化——并行执行

- 并行实行：
  - 同步执行hive的多个阶段，hive在执行过程，将一个查询转化成一个或者多个阶段。某个特定的job可能包含众多的阶段，而这些阶段可能并非完全相互依赖的，也就是说可以并行执行的，这样可能使得整个job的执行时间缩短。hive执行开启：set hive.exec.parallel=true



## Hive 的优化——数据倾斜

### • 操作

- Join
- Group by
- Count Distinct

### • 原因

- key分布不均导致的
- 人为的建表疏忽
- 业务数据特点

### • 症状

- 任务进度长时间维持在99%（或100%），查看任务监控页面，发现只有少量（1个或几个）reduce子任务未完成。
- 查看未完成的子任务，可以看到本地读写数据量积累非常大，通常超过10GB可以认定为发生数据倾斜。

### • 倾斜度

- 平均记录数超过50w且最大记录数是超过平均记录数的4倍。
- 最长时长比平均时长超过4分钟，且最大时长超过平均时长的2倍。

### • 万能方法

- `hive.groupby.skewindata=true`

## Hive 的优化——数据倾斜——大小表关联

### • 原因

- Hive在进行join时，按照join的key进行分发，而在join左边的表的数据会首先读入内存，如果左边表的key相对分散，读入内存的数据会比较小，join任务执行会比较快；而如果左边的表key比较集中，而这张表的数据量很大，那么数据倾斜就会比较严重，而如果这张表是小表，则还是应该把这张表放在join左边。

### • 思路

- 将key相对分散，并且数据量小的表放在join的左边，这样可以有效减少内存溢出错误发生的几率
- 使用map join让小的维度表先进内存。

### • 方法

- Small\_table join big\_table

## Hive 的优化——数据倾斜——大大表关联

- 原因

- 日志中有一部分的userid是空或者是0的情况，导致在用user\_id进行hash分桶的时候，会将日志中userid为0或者空的数据分到一起，导致了过大的斜率。

- 思路

- 把空值的key变成一个字符串加上随机数，把倾斜的数据分到不同的reduce上，由于null值关联不上，处理后并不影响最终结果。

- 方法

- `on case when (x.uid = '-' or x.uid = '0 ' or x.uid is null) then concat('dp_hive_search',rand()) else x.uid  
end = f.user_id;`

## Hive 的优化——数据倾斜——大表关联（业务削减）

## • 案例

- `Select * from dw_log t join dw_user t1 on t.user_id=t1.user_id`
- 现象：两个表都上千万，跑起来很悬

## • 思路

- 当天登陆的用户其实很少

## • 方法

- `Select/*+MAPJOIN(t12)*/ *`
- `from dw_log t1`
- `join (`
- `select/*+MAPJOIN(t)*/ t1.*`
- `from (`
- `select user_id from dw_log group by user_id`
- `) t`
- `join dw_user t1`
- `on t.user_id=t1.user_id`
- `) t12`
- `on t11.user_id=t12.user_id`

八斗大数据内部资料，盗版必究——

## Hive 的优化——数据倾斜——聚合时存在大量特殊值

### • 原因

- 做count distinct时，该字段存在大量值为NULL或空的记录。

### • 思路

- count distinct时，将值为空的情况单独处理，如果是计算count distinct，可以不用处理，直接过滤，在最后结果中加1。
- 如果还有其他计算，需要进行group by，可以先将值为空的记录单独处理，再和其他计算结果进行union

### • 方法

- `select cast(count(distinct(user_id))+1 as bigint) as user_cnt`
- `from tab_a`
- `where user_id is not null and user_id <> ''`

## Hive 的优化——数据倾斜——空间换时间

## • 案例

- Select day,count(distinct session\_id),count(distinct user\_id) from log a group by day

## • 问题

- 同一个reduce上进行distinct操作时压力很大

## • 方法

- select day,
- count(case when type='session' then 1 else null end) as session\_cnt,
- count(case when type='user' then 1 else null end) as user\_cnt
- from (
- select day,session\_id,type
- from (
- select day,session\_id,'session' as type
- from log
- union all
- select day user\_id,'user' as type
- from log
- )
- group by day,session\_id,type
- ) t1
- group by day



OutLine

Hive基础

【实践】Hive搭建

【实践】Hive练习

## M y s q l 配 置

- 默认情况下，Hive的元数据信息存储在内置的Derby数据中。
- Hive支持将元数据存储在MySQL中
- 元数据存储配置：
  - 【本地配置1】：默认
  - 【本地配置2】：本地搭建mysql，通过localhost:Port方式访问
  - 【远程配置】：远程搭建mysql，通过IP:Port方式访问

## MySQL 配置

- **第一步：** 安装MySQL服务器端和MySQL客户端,并启动MySQL服务
- 安装：
  - yum install mysql
  - yum install mysql-server
- 启动：
  - /etc/init.d/mysqld start
- 设置用户名和密码mysqladmin -u root password码：
  - '111111 '
- 测试登录是否成功：
  - mysql -uroot -p111111

## M y s q l 配 置

- **第二步：安装Hive**
- 下载apache-hive-0.13.0-bin.tgz，并解压：
- 在conf目录下，创建hive-site.xml配置文件：

```
[root@master conf]# cat hive-site.xml
<configuration>
  <property>
    <name>javax.jdo.option.ConnectionURL</name>
    <value>jdbc:mysql://localhost:3306/hive?createDatabaseIfNotExist=true</value>
  </property>

  <property>
    <name>javax.jdo.option.ConnectionDriverName</name>
    <value>com.mysql.jdbc.Driver</value>
  </property>

  <property>
    <name>javax.jdo.option.ConnectionUserName</name>
    <value>root</value>
  </property>

  <property>
    <name>javax.jdo.option.ConnectionPassword</name>
    <value>111111</value>
  </property>
</configuration>
```

## M y s q l 配 置

- **第二步：**安装Hive
- 修改bashrc，配置环境变量：

```
# hive conf  
export HIVE_HOME=/usr/local/src/hive-0.12.0-bin  
export PATH=$MAHOUT_HOME/conf:$MAHOUT_HOME/bin:$ZOOKEEPER_HOME/bin:$HIVE_HOME/bin:$PATH
```

- 将mysql-connector-java-5.1.41-bin.jar拷贝到hive home的lib目录下，以支持hive对mysql的操作

## MySQL 配置

- 第三步：测试Hive

```
[root@master badou]# hive
Logging initialized using configuration in jar:file:/usr/lo
hive> show tables;
OK
Time taken: 5.313 seconds
hive> create EXTERNAL TABLE w_a
> (
> usrid STRING,
> age STRING,
> sex STRING
> )
> ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
>                                LINES TERMINATED BY '\n'
> ;
OK
Time taken: 0.803 seconds
hive> show tables;
OK
w_a
Time taken: 0.052 seconds, Fetched: 1 row(s)
hive> desc w_a;
OK
usrid      string      None
age        string      None
sex        string      None
Time taken: 0.195 seconds, Fetched: 3 row(s)
hive> select * from w_a;
OK
user1  27      1
user2  28      1
user3  29      0
user4  30      1
user5  31      0
user6  32      1
user7  33      1
user8  34      0
Time taken: 0.343 seconds, Fetched: 8 row(s)
hive>
```

```
mysql> show tables;
+-----+
| Tables_in_hive |
+-----+
| BUCKETING_COLS |
| CDS              |
| COLUMNS_V2      |
| DATABASE_PARAMS  |
| DBS              |
| IDX5             |
| INDEX_PARAMS     |
| PARTITIONS       |
| PARTITION_KEYS   |
| PARTITION_KEY_VALS |
| PARTITION_PARAMS |
| PART_COL_PRIVS   |
| PART_PRIVS       |
| SDS              |
| SD_PARAMS        |
| SEQUENCE_TABLE   |
| SERDES           |
| SERDE_PARAMS     |
| SKEWED_COL_NAMES |
| SKEWED_COL_VALUE_LOC_MAP |
| SKEWED_STRING_LIST |
| SKEWED_STRING_LIST_VALUES |
| SKEWED_VALUES    |
| SORT_COLS        |
| TABLE_PARAMS    |
| TAB_COL_STATS    |
| TBLS             |
| TBL_COL_PRIVS    |
| TBL_PRIVS        |
| VERSION           |
+-----+
30 rows in set (0.00 sec)
```

OutLine

Hive基础

【实践】Hive搭建

【实践】Hive练习

## 案例一：导入本地 Local 的数据，并进行简单统计

- 准备数据
- 设计schema, 建库, 建表

```
[root@master hive_test]# vim 'create_table.sql'
1 create EXTERNAL TABLE w_a
2 (
3     usrid STRING,
4     age STRING,
5     sex STRING
6 )
7 ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
8   LINES TERMINATED BY '\n';
9
10 create EXTERNAL TABLE w_b
11 (
12     usrid STRING,
13     active STRING,
14     time STRING
15 )
16 ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
17   LINES TERMINATED BY '\n'
```

```
[root@master badou]# hadoop fs -ls /user/hive/warehouse
Found 2 items
drwxr-xr-x  - root supergroup          0 2017-03-27 01:19 /user/hive/warehouse/w_a
drwxr-xr-x  - root supergroup          0 2017-03-27 01:19 /user/hive/warehouse/w_b
```



## 案例一：导入本地 Local 的数据，并进行简单统计

## • 导入

```
[root@master hive_test]# cat insert.sh
hive -e "LOAD DATA LOCAL INPATH '/home/badou/hive_test/a.txt' OVERWRITE INTO TABLE w_a"
hive -e "LOAD DATA LOCAL INPATH '/home/badou/hive_test/b.txt' OVERWRITE INTO TABLE w_b"
[root@master hive_test]#
```

```
[root@master hive_test]# cat a.txt
user1 27 1
user2 28 1
user3 29 0
user4 30 1
user5 31 0
user6 32 1
user7 33 1
user8 34 0
```



```
hive> select * from w_a;
OK
user1 27 1
user2 28 1
user3 29 0
user4 30 1
user5 31 0
user6 32 1
user7 33 1
user8 34 0
Time taken: 0.077 seconds, Fetched: 8 row(s)
```

```
[root@master badou]# hadoop fs -ls /user/hive/warehouse/w_a
Found 1 items
-rw-r--r-- 3 root supergroup 88 2017-03-27 01:27 /user/hive/warehouse/w_a/a.txt
[root@master badou]# hadoop fs -ls /user/hive/warehouse/w_b
Found 1 items
-rw-r--r-- 3 root supergroup 114 2017-03-27 01:27 /user/hive/warehouse/w_b/b.txt
[root@master badou]#
```

## 案例二：两表的Join

- 执行命令：`select a.*, b.* from w_a a join w_b b on a.usrid=b.usrid;`

```
hive> select a.*, b.* from w_a a join w_b b on a.usrid=b.usrid;
Total MapReduce jobs = 1
Execution log at: /tmp/root/.log
2017-03-27 01:30:07 Starting to launch local task to process map join; maximum me
2017-03-27 01:30:08 Dump the side-table into file: file:/tmp/root/hive_2017-03-27_01-2
00--.hashtable
2017-03-27 01:30:08 Upload 1 File to: file:/tmp/root/hive_2017-03-27_01-29-54_890_2166
e
2017-03-27 01:30:08 End of local task; Time Taken: 1.242 sec.
Execution completed successfully
Mapred Local Task Succeeded . Convert the Join into MapJoin
Mapred Local Task Succeeded . Convert the Join into MapJoin
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_201703260516_0001, Tracking URL = http://master:50030/jobdetails.jsp?j
Kill Command = /usr/local/src/hadoop-1.2.1/libexec/./bin/hadoop job -kill job_2017032605
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2017-03-27 01:30:21,296 Stage-3 map = 0%, reduce = 0%
2017-03-27 01:30:29,379 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.77 sec
2017-03-27 01:30:30,386 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.77 sec
2017-03-27 01:30:31,395 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.77 sec
2017-03-27 01:30:32,402 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.77 sec
2017-03-27 01:30:33,411 Stage-3 map = 100%, reduce = 100%, Cumulative CPU 1.77 sec
MapReduce Total cumulative CPU time: 1 seconds 770 msec
Ended Job = job_201703260516_0001
MapReduce Jobs Launched:
Job 0: Map: 1 Cumulative CPU: 1.77 sec HDFS Read: 324 HDFS Write: 180 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 770 msec
OK
user1 27 1 user1 100 20170301
user3 29 0 user3 101 20170302
user4 30 1 user4 102 20170303
user5 31 0 user5 103 20170304
user7 33 1 user7 104 20170305
user8 34 0 user8 105 20170306
Time taken: 38.673 seconds, Fetched: 6 row(s)
hive>
```

MapReduce!!!

## 案例三：UDF

- UDF: User-Defined-Function 用户自定义函数
- UDF函数可以直接应用于select语句，对查询结构做格式化处理后，再输出内容。
- 编写UDF函数的时候需要注意以下几点：
  - 自定义UDF需要继承org.apache.hadoop.hive.ql.UDF。
  - 需要实现evaluate函。
  - evaluate函数支持重载。

## 案例三：UDF

- 借助UDF，实现一个大写函数
- 效果验证：

```
hive> select userid,uppercase(usrid),age from w_a;  
FAILED: SemanticException [Error 10004]: Line 1:7 Invalid table alias or  
hive> select usrid,uppercase(usrid),age from w_a;  
Total MapReduce jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks is set to 0 since there's no reduce operator  
Starting Job = job_201703260516_0004, Tracking URL = http://master:50030  
Kill Command = /usr/local/src/hadoop-1.2.1/libexec/./bin/hadoop job -k  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1  
2017-03-27 02:17:35,299 Stage-1 map = 0%, reduce = 0%  
2017-03-27 02:17:40,325 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.55 sec  
2017-03-27 02:17:41,331 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.55 sec  
2017-03-27 02:17:42,336 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.55 sec  
2017-03-27 02:17:43,341 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.55 sec  
2017-03-27 02:17:44,348 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 1.55 sec  
MapReduce Total cumulative CPU time: 1 seconds 550 msec  
Ended Job = job_201703260516_0004  
MapReduce Jobs Launched:  
Job 0: Map: 1 Cumulative CPU: 1.55 sec HDFS Read: 298 HDFS Write: 120 SUCCESS  
Total MapReduce CPU Time Spent: 1 seconds 550 msec  
OK  
user1  USER1  27  
user2  USER2  28  
user3  USER3  29  
user4  USER4  30  
user5  USER5  31  
user6  USER6  32  
user7  USER7  33  
user8  USER8  34  
Time taken: 18.245 seconds, Fetched: 8 row(s)  
hive>
```

Uppercase.java

```
1 package com.badou.hive.udf;  
2  
3 import org.apache.hadoop.hive.ql.exec.UDF;  
4 import org.apache.hadoop.io.Text;  
5  
6 public class Uppercase extends UDF {  
7  
8     public Text evaluate(final Text s) {  
9         return new Text(s.toString().toUpperCase());  
10    }  
11 }  
12
```

## 案例四：从HDFS中导入

- 执行命令：
  - `LOAD DATA INPATH '/user_name.data.utf.txt' OVERWRITE INTO TABLE u_info`
- `overwrite`表示加载的数据会覆盖原来的内容
- 对比本地的方式：`LOAD DATA LOCAL INPATH`

## 案例五：利用 Insert 命令导入数据

- 执行命令（例子）：

insert into table table1 select usrid, age from w\_a limit 3;

- 也可以支持动态分区插入：

insert into table test1 partition(c) select \* from test2;

```
hive> create table table1(a string, b string);
OK
Time taken: 0.037 seconds
hive> insert into table table1 select usrid, age from w_a limit 3;
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_201703260516_0005, Tracking URL = http://master:50030/jobdetail
Kill Command = /usr/local/src/hadoop-1.2.1/libexec/./bin/hadoop job -kill job_
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2017-03-27 04:24:05,185 Stage-1 map = 0%, reduce = 0%
2017-03-27 04:24:10,218 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:11,224 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:12,230 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:13,236 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:14,241 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:15,246 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:16,251 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:17,257 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:18,265 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:19,272 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec
2017-03-27 04:24:20,281 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.64 sec
2017-03-27 04:24:21,288 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.64 sec
2017-03-27 04:24:22,295 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.64 sec
2017-03-27 04:24:23,302 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.64 sec
MapReduce Total cumulative CPU time: 3 seconds 640 msec
Ended Job = job_201703260516_0005
Loading data to table default.table1
Table default.table1 stats: [num_partitions: 0, num_files: 1, num_rows: 0, total
MapReduce Jobs Launched:
Job 0: Map: 1 Reduce: 1 Cumulative CPU: 3.64 sec HDFS Read: 298 HDFS Write:
Total MapReduce CPU Time Spent: 3 seconds 640 msec
OK
Time taken: 29.245 seconds
hive> select * from table1;
OK
user1 27
user2 28
user3 29
Time taken: 0.066 seconds, Fetched: 3 row(s)
hive>
```



## 案例六：直接通过查询插入

- 执行命令（例子）：
- `create table test2 as select * from test1;`

```
hive> create table test1 as select * from w_b;
Total MapReduce jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no red
Starting Job = job_201703260516_0006, Tracking URL = ht
Kill Command = /usr/local/src/hadoop-1.2.1/libexec/..b
Hadoop job information for Stage-1: number of mappers:
2017-03-27 04:32:06,579 Stage-1 map = 0%, reduce = 0%
2017-03-27 04:32:10,601 Stage-1 map = 100%, reduce = 0%
2017-03-27 04:32:11,607 Stage-1 map = 100%, reduce = 0%
2017-03-27 04:32:12,613 Stage-1 map = 100%, reduce = 0%
2017-03-27 04:32:13,622 Stage-1 map = 100%, reduce = 0%
2017-03-27 04:32:14,631 Stage-1 map = 100%, reduce = 1
MapReduce Total cumulative CPU time: 890 msec
Ended Job = job_201703260516_0006
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://192.168.183.10:9000/tmp/hive-roo
Moving data to: hdfs://192.168.183.10:9000/user/hive/wa
Table default.test1 stats: [num_partitions: 0, num_file
MapReduce Jobs Launched:
Job 0: Map: 1 Cumulative CPU: 0.89 sec HDFS Read: 3
Total MapReduce CPU Time Spent: 890 msec
OK
Time taken: 14.636 seconds
hive> select * from test1;
OK
user1 100 20170301
user3 101 20170302
user4 102 20170303
user5 103 20170304
user7 104 20170305
user8 105 20170306
Time taken: 0.07 seconds, Fetched: 6 row(s)
hive> select * from w_b;
OK
user1 100 20170301
user3 101 20170302
user4 102 20170303
user5 103 20170304
user7 104 20170305
user8 105 20170306
Time taken: 0.075 seconds, Fetched: 6 row(s)
hive>
```

## 案例六：数据导出（导出为本地文件）

- 执行命令（例子）：
- insert overwrite local directory  
'/home/badou/hive\_test/1.txt' select usrid,  
sex from w\_a;

```
[root@master hive_test]# ls 1.txt/  
000000_0  
[root@master hive_test]# cat 1.txt/000000_0  
user101  
user201  
user300  
user401  
user500  
user601  
user701  
user800
```

```
hive> insert overwrite local directory '/home/badou/hive_test/1.txt' select usrid, sex from w_a;  
Total MapReduce jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks is set to 0 since there's no reduce operator  
Starting Job = job_201703260516_0007, Tracking URL = http://master:50030/jobdetails.jsp?jobid=job_201703260516_0007  
Kill Command = /usr/local/src/hadoop-1.2.1/libexec/./bin/hadoop job -kill job_201703260516_0007  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0  
2017-03-27 04:36:45,821 Stage-1 map = 0%, reduce = 0%  
2017-03-27 04:36:49,841 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 0.96 sec  
2017-03-27 04:36:50,848 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 0.96 sec  
2017-03-27 04:36:51,853 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 0.96 sec  
2017-03-27 04:36:52,858 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 0.96 sec  
MapReduce Total cumulative CPU time: 960 msec  
Ended Job = job_201703260516_0007  
Copying data to local directory /home/badou/hive_test/1.txt  
Copying data to local directory /home/badou/hive_test/1.txt  
MapReduce Jobs Launched:  
Job 0: Map: 1 Cumulative CPU: 0.96 sec HDFS Read: 298 HDFS Write: 64 SUCCESS  
Total MapReduce CPU Time Spent: 960 msec  
OK  
Time taken: 12.816 seconds  
hive>
```



## 案例七：数据导出（导出为HDFS文件）

- 执行命令（例子）：insert overwrite directory '/hive\_output' select \* from w\_b;

```
hive> insert overwrite directory '/hive_output' select * from w_b;
Total MapReduce jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_201703260516_0008, Tracking URL = http://master:50030/jobdetails
Kill Command = /usr/local/src/hadoop-1.2.1/libexec/./bin/hadoop job -kill job_201703260516_0008
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2017-03-27 04:41:07,956 Stage-1 map = 0%, reduce = 0%
2017-03-27 04:41:11,978 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.04 sec
2017-03-27 04:41:12,988 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.04 sec
2017-03-27 04:41:13,994 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.04 sec
2017-03-27 04:41:15,004 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 1.04 sec
MapReduce Total cumulative CPU time: 1 seconds 40 msec
Ended Job = job_201703260516_0008
Stage-3 is selected by condition resolver.
Stage-2 is filtered out by condition resolver.
Stage-4 is filtered out by condition resolver.
Moving data to: hdfs://192.168.183.10:9000/tmp/hive-root/hive_2017-03-27_04-41-15
Moving data to: /hive_output
MapReduce Jobs Launched:
Job 0: Map: 1 Cumulative CPU: 1.04 sec HDFS Read: 32 KB HDFS Write: 114 KB
Total MapReduce CPU Time Spent: 1 seconds 40 msec
OK
Time taken: 12.907 seconds
hive>
```

```
[root@master badou]# hadoop fs -ls /hive_output
Found 1 items
-rw-r--r-- 3 root supergroup 114 2017-03-27 04:41 /hive_output/000000_0
[root@master badou]# hadoop fs -text /hive_output/000000_0
user10020170301
user3010120170302
user4010220170303
user5010320170304
user7010420170305
user8010520170306
[root@master badou]#
```

## 案例八：Partition

- partition是Hive提供的一种机制：用户通过指定一个或多个partition key，决定数据存放方式，进而优化数据的查询,一个表可以指定多个partition key，每个partition在hive中以文件夹的形式存在。

```
hive> LOAD DATA LOCAL INPATH '/home/badou/hive_test/p1.txt' OVERWRITE INTO TABLE p_t partition(dt='20170302');
Copying data from file:/home/badou/hive_test/p1.txt
Copying file: file:/home/badou/hive_test/p1.txt
Loading data to table default.p_t partition (dt=20170302)
Partition default.p_t{dt=20170302} stats: [num_files: 1, num_rows: 0, total_size: 72, raw_data_size: 0]
Table default.p_t stats: [num_partitions: 3, num_files: 3, num_rows: 0, total_size: 360, raw_data_size: 0]
OK
Time taken: 0.699 seconds
hive> LOAD DATA LOCAL INPATH '/home/badou/hive_test/p2.txt' OVERWRITE INTO TABLE p_t partition(dt='20170303');
Copying data from file:/home/badou/hive_test/p2.txt
Copying file: file:/home/badou/hive_test/p2.txt
Loading data to table default.p_t partition (dt=20170303)
Partition default.p_t{dt=20170303} stats: [num_files: 1, num_rows: 0, total_size: 36, raw_data_size: 0]
Table default.p_t stats: [num_partitions: 4, num_files: 4, num_rows: 0, total_size: 396, raw_data_size: 0]
OK
Time taken: 0.54 seconds
hive> select * from p_t;
OK
user2  28      20170302
user4  30      20170302
user6  32      20170302
user8  34      20170302
user1  27      20170303
user7  33      20170303
Time taken: 0.087 seconds, Fetched: 6 row(s)
hive> select * from p_t where dt='20170303';
OK
user1  27      20170303
user7  33      20170303
Time taken: 0.171 seconds, Fetched: 2 row(s)
hive>
```

```
[root@master badou]# hadoop fs -ls /user/hive/warehouse/p_t
Found 2 items
drwxr-xr-x - root supergroup 0 2017-03-27 05:01 /user/hive/warehouse/p_t/dt=20170302
drwxr-xr-x - root supergroup 0 2017-03-27 05:01 /user/hive/warehouse/p_t/dt=20170303
[root@master badou]#
```

```
[root@master hive_test]# vim create_partition.sql
[root@master hive_test]# vim create_partition.sql
1 create TABLE p_t
2 (
3     usrid STRING,
4     age STRING
5 )
6 PARTITIONED BY (dt STRING)
7 ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
8 LINES TERMINATED BY '\n';
```

```
hive> show tables;
OK
p_t
table1
test1
w_a
w_b
Time taken: 0.023 seconds, Fetched: 5 row(s)
hive> desc p_t;
OK
usrid      string      None
age        string      None
dt         string      None

# Partition Information
# col_name  data_type   comment
dt          string      None
Time taken: 0.172 seconds, Fetched: 8 row(s)
hive>
```

```
[root@master hive_test]# cat p1.txt
user2  28      20170301
user4  30      20170301
user6  32      20170301
user8  34      20170301
[root@master hive_test]# cat p2.txt
user1  27      20170303
user7  33      20170303
[root@master hive_test]#
```

## 案例九：Transform

- transform功能部分可以用UDF替代，但是如果拼接的字段是根据上一次查询的结果时，UDF就不能用，UDF只能用在本行操作
- transform功能缺点是效率底了点

官网的一个case:


```
FROM (
    FROM pv_users
    SELECT TRANSFORM(pv_users.userid, pv_users.date)
    USING 'map_script'
    AS dt, uid
    CLUSTER BY dt
) map_output
INSERT OVERWRITE TABLE pv_users_reduced
SELECT TRANSFORM(map_output.dt, map_output.uid)
USING 'reduce_script'
AS date, count;
```

```
[root@master hive_test]# cat transform.awk
{
    print $1"_"$2
}
```

```
hive> ADD FILE /home/badou/hive_test/transform.awk;
Added resource: /home/badou/hive_test/transform.awk
```


```
hive> SELECT TRANSFORM(userid, age) using "awk -f transform.awk" as (uuu) from w_a;
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_201703260516_0012, Tracking URL = http://master:50030/jobdetails
Kill Command = /usr/local/src/hadoop-1.2.1/libexec/..bin/hadoop job -kill job_201
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2017-03-27 05:30:01,893 Stage-1 map = 0%, reduce = 0%
2017-03-27 05:30:05,916 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.06 sec
2017-03-27 05:30:06,922 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.06 sec
2017-03-27 05:30:07,931 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.06 sec
2017-03-27 05:30:08,936 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 1.06 sec
MapReduce Total cumulative CPU time: 1 seconds 60 msec
Ended Job = job_201703260516_0012
MapReduce Jobs Launched:
Job 0: Map: 1 Cumulative CPU: 1.06 sec HDFS Read: 298 HDFS Write: 72 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 60 msec
OK
user1_27
user2_28
user3_29
user4_30
user5_31
user6_32
user7_33
user8_34
Time taken: 12.991 seconds, Fetched: 8 row(s)
hive>
```

## 案例十：Hive 整合 Hbase



## 案例十：Hive整合Hbase

- 创建Hbase表：
  - create 'classes','user'
- 加入数据：
  - put 'classes','001','user:name','jack'
  - put 'classes','001','user:age','20'
  - put 'classes','002','user:name','liza'
  - put 'classes','002','user:age','18'





## 案例十：Hive整合Hbase

- 创建Hive表并验证：
  - create external table classes(id int, name string, age int)
  - STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
  - WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,user:name,user:age")
  - TBLPROPERTIES("hbase.table.name" = "classes");
- 再添加数据到Hbase：
  - put 'classes','003','user:age','1820183291839132'

---

# Q & A

@八斗学院

---